

# The Gretl `fsreg` function package for running the forward stagewise regression

Artur Tarassow

Version 0.1

## Changelog

- Version 0.1 (August, 2020)

— initial release

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>The FSREG algorithm</b>	<b>2</b>
<b>3</b>	<b>Install and load the package</b>	<b>3</b>
<b>4</b>	<b>Example</b>	<b>4</b>
<b>5</b>	<b>Using the GUI</b>	<b>6</b>
<b>6</b>	<b>Public functions and parameter values</b>	<b>6</b>
6.1	<code>fsreg()</code> . . . . .	7
6.2	<code>print_fsreg_results()</code> . . . . .	7
6.3	<code>plot_rho_values()</code> . . . . .	7
6.4	<code>plot_coefficient_paths()</code> . . . . .	8

## 1 Introduction

Selecting the relevant predictors among a many potential ones is a crucial task. Neglecting relevant predictors may lead to inconsistent parameter estimates while considering irrelevant regressors yields an inefficient estimator. Furthermore, considering highly correlated predictors into a standard least square regression most probably suffers from multicollinearity issues. Lastly, standard least squares cannot handle the case when the number of observations,  $T$ , exceeds the number of potential regressors,  $k$ .

So called shrinkage and/ or selection estimators such as Ridge or Lasso among others, are known to handle such issues. Another estimation approach marks the so called forward stagewise regression approach (FSREG henceforth). FSREG follows a simple strategy for constructing a sequence of sparse regression estimates: Initially set all coefficients to zero, and iteratively update the coefficient (by a small amount, depending on the *learning rate*  $\epsilon$ ) of the variable that achieves the maximal absolute correlation with the current residual. *Learning* from the residuals has some connection to Boosting. However, this procedure also has some interesting connection to the Lasso under some conditions (REFs). As the step size goes to zero, the sequence of forward stagewise estimates exactly coincides with the lasso path. While, this equivalence holds outside of least squares regression, currently we only support minimization of squared loss (RMSE).

## 2 The FSREG algorithm

The FSREG algorithm works as follows:<sup>1</sup>

1. Start with  $r = y$  and  $\beta_1 = \beta_2 = \dots = \beta_k = 0$ .
2. Find the predictor  $x_j$  most correlated with  $r$ .
3. Update the  $j$ -th predictor  $\beta_j^i \leftarrow \beta_j^{i-1} + \delta_j$  where  $\delta_j = \epsilon \times \text{sign} \langle r^{i-1}, x_j \rangle$
4. Update the residuals  $r^i \leftarrow r^{i-1} - \delta_j \times x_j$  and repeat steps 2 and 3 many times.

Here  $y$ ,  $\beta$ ,  $\epsilon$ ,  $\langle r, x_j \rangle$  and  $i$  refer to the endogenous variable, the unknown regression coefficients, the learning rate, the correlation between the current residuals and the  $j$ -th regressor and the  $i$ -th iteration. The learning rate  $\epsilon$  is a hyper-parameter and set to a fixed constant (e.g.,  $\epsilon = 0.01$ ). The only computations intense task is to compute the correlation between  $r$  and the all  $k$  predictors. We make use of Gretl's `mcorr()` function for this which is written in C and hence computation is very fast.

The idea behind the stagewise updates is simple: at each iteration greedily select the predictor  $j$  that has the largest absolute inner product (or correlation, for standardized variables) with the residual. As the residuals refer to the yet unexplained part of the model, we are searching for any variable that still has some information content for explaining 'something' left unexplained.

Given the greediness that only a single predictor is selected each iteration, updating the coefficient of variable  $j$  by a large amount is problematic. Instead, the parameter  $\epsilon$  slows down the learning process only changing the coefficient by a tiny amount. Thus, many iterations are required to yield reasonable parameter estimates among a large set of potential predictor variables. Figure 1 illustrates the coefficient path of the coefficient estimates over 2000 iterations.

Early stopping rules are important for two reasons: First, one tries to avoid over-fitting, and second it may be unnecessary to run  $\bar{N}$  iterations if no improvement (in terms of model fit) can be seen after  $N \ll \bar{N}$  iterations. The implemented early stopping strategy checks the absolute correlation  $|\langle r, x_j^i \rangle|$ : In case the absolute correlation does not improve for  $n$  (e.g.,  $n = 30$ )

---

<sup>1</sup>Also note, that this is analogous to least squares boosting, with the number of trees equal to the number of predictors predictors.




Figure 1: Coefficient paths

iterations, we assume that the coefficient estimates have converged and stop the algorithm. Figure 2 illustrates the development of the development of absolute correlation between the residuals and the selected variables.




Figure 2: Correlation dynamics

### 3 Install and load the package

The `fsreg` package is publicly available on the gretl server. The package must be downloaded once, and loaded into memory each time gretl is started.

```
clear
set verbose off

pkg install fsreg      # Download package (only once need)
include fsreg.gfn      # Load the package into memory
help fsreg             # Open the help file
```

## 4 Example

For illustration we use the “mroz” cross-sectional data set comprising 753 observations from a study of income dynamics. The endogenous variable is named **wage** and refers to the a wife’s 1975 average hourly earnings (in 1975 dollars). The data set includes additional 24 potential predictors.

The sample script opens the data set, sets the right-hand-side list of predictors and computes standard least square estimates.

```
open mroz.gdt --quiet
list RHS = const dataset
RHS -= wage      # drop lhs variable
ols LHS RHS      # run ols as benchmark
```

The OLS output is:

m: OLS, using observations 1-753

Dependent variable: LHS

	coefficient	std. error	t-ratio	p-value	
-----	-----	-----	-----	-----	-----
const	4.10278	2.37643	1.726	0.0847	*
taxableinc	1.31268e-05	2.74615e-05	0.4780	0.6328	
federaltax	2.12780e-05	9.20550e-05	0.2311	0.8173	
hsiblings	0.000386927	0.0355359	0.01089	0.9913	
hfathereduc	0.0139611	0.0285757	0.4886	0.6253	
hmothereduc	0.0396817	0.0276745	1.434	0.1520	
siblings	0.0694937	0.0358372	1.939	0.0529	*
lfp	3.45220	0.263461	13.10	2.25e-35	***
hours	0.00111994	0.000152678	7.335	5.91e-13	***
kids16	0.0423818	0.179218	0.2365	0.8131	
kids618	0.0329783	0.0703111	0.4690	0.6392	
age	0.00589606	0.0226630	0.2602	0.7948	
educ	0.216125	0.0492212	4.391	1.30e-05	***
wage76	0.543352	0.0463328	11.73	3.26e-29	***
hhours	0.000484660	0.000170424	2.844	0.0046	***
hage	0.00327033	0.0217384	0.1504	0.8805	
heduc	0.0511196	0.0359926	1.420	0.1560	
hwage	0.123328	0.0380489	3.241	0.0012	***
faminc	3.36999e-05	1.52480e-05	2.210	0.0274	**
mtr	4.76933	2.38306	2.001	0.0457	**
mothereduc	0.0282278	0.0299406	0.9428	0.3461	
fathereduc	0.0220970	0.0285042	0.7752	0.4385	
unemployment	0.0156837	0.0262039	0.5985	0.5497	
largecity	0.00659037	0.179593	0.03670	0.9707	
exper	0.00247946	0.0121007	0.2049	0.8377	
Mean dependent var	2.374565	S.D. dependent var	3.241829		
Sum squared resid	3351.283	S.E. of regression	2.145556		
R-squared	0.575954	Adjusted R-squared	0.561974		
F(24, 728)	41.19977	P-value(F)	2.5e-118		
Log-likelihood	1630.588	Akaike criterion	3311.176		
Schwarz criterion	3426.777	Hannan-Quinn	3355.711		

Excluding the constant, p-value was highest for variable 3 (hsiblings)

Next, we run the forward stagewise regression by calling the `fsreg()` function:

```
bundle B = fsreg(LHS, RHS)
print_fsreg_results(B)      # Print estimation results
```

The regression results are as follows (note that inference is not supported, yet!):

```
Info: No improvement in correlation for the last 50 iterations.
Early stopping applies.
```

```
Forward-stagewise regression results (no inference)
```

```
-----
```

	coefficient	std. error	z	p-value
-----	-----	-----	-----	-----
const	2.88386	NA	NA	NA
taxableinc	-1.36310e-05	NA	NA	NA
hmothereeduc	-0.0483092	NA	NA	NA
siblings	-0.0700839	NA	NA	NA
lfp	3.27042	NA	NA	NA
hours	-0.000930155	NA	NA	NA
educ	0.213255	NA	NA	NA
wage76	0.535865	NA	NA	NA
hhours	-0.000272163	NA	NA	NA
heduc	-0.0536584	NA	NA	NA
hwage	-0.0766289	NA	NA	NA
faminc	2.65937e-05	NA	NA	NA
mtr	-3.88264	NA	NA	NA
mothereeduc	-0.0481345	NA	NA	NA

```
Number of iterations = 2068
Correl. w. residuals = -0.0329421
S.E. of regression = 2.12071
R-squared = 0.572159
```

The estimator converged after 2068 iterations and selected only 14 out of 25 potential predictors. The last correlation coefficient between the final residuals and some predictor was found being close to zero (-0.033). Even though only 14 predictors are selected being relevant, the  $R^2$  statistics is of similar size of the OLS equivalent. Also the standard error of the residuals is slightly smaller with 2.12 compared to the OLS-based value of 2.14.

## 5 Using the GUI

Instead of scripting, one may access the `fsreg` procedure by means of the Gretl GUI. Once the package is installed and loaded, simply open the menu “*Model -> Other linear models -> Forward Stagewise*”. This a menu window as depicted in Figure 3.

## 6 Public functions and parameter values

The following public functions currently exist.

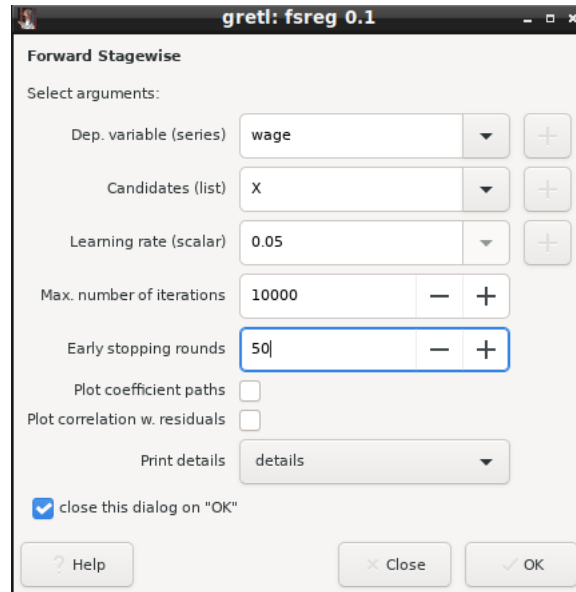


Figure 3: GUI access window

## 6.1 fsreg()

The `fsreg()` function marks the main function. The function arguments are:

```
fsreg(const series y, const list X, bundle opts[null])
```

Return type: `bundle`

Argument	Description
<code>y</code>	series, Endogenous variable
<code>X</code>	list, Non-empty list of predictor variables
<code>opts</code>	bundle, Pass parameters for controlling the algorithm (optional)

Return type: `bundle`

The returned bundle includes various which are listed in Table 1.

The additional parameters which can be passed by means of the `opts` bundle to `fsreg()` are shown in Table 2.

## 6.2 print\_fsreg\_results()

The `print_fsreg_results()` function takes the resulting bundle returned by the `fsreg()` function, and prints a summary of the estimation results. The argument is:

```
print_fsreg_results(const bundle B)
```

Return type: `void`

## 6.3 plot\_rho\_values()

For plotting the development of the remaining correlation with the residuals, simply call the `plot_rho_values()` function. It takes the resulting bundle returned by the `fsreg()` function. The argument is:

Parameter	Description
<code>rho_values</code>	Vector holding correlation coefficients with the residuals for each iteration
<code>max_num_iterations</code>	Number of the maximum iterations
<code>actual_num_iterations</code>	Actual number of iterations ran
<code>eta</code>	Learning rate
<code>early_stopping_rounds</code>	Number of iterations of no improvement before stopping
<code>yname</code>	String holding the name of the endogenous variable
<code>Xnames_wo_constant</code>	String array holding the names of all predictor variables without the constant
<code>Xnames</code>	String array holding the names of all predictor variables incl. the constant
<code>X_final</code>	List incl. finally selected predictors.
<code>betas</code>	Matrix holding the coefficient point estimates across iterations (rows) for each predictor (columns)
<code>coeff_nonzero</code>	k by 1 vector incl. the final coefficient point estimates for all selected predictors (non-zero coefficients)
<code>coeff</code>	n by 1 vector incl. the coefficient point estimates of all predictors (incl. zero coefficients)
<code>with_constant</code>	Boolean taking zero if the passed list X does not incl. an intercept, otherwise one
<code>verbose</code>	Integer indicating the level of verbosity
<code>T</code>	Number of effective (non-missing) observations.
<code>yhat</code>	Fitted values using final point coefficient estimates
<code>uhat</code>	Estimated final residuals
<code>uhat_variance</code>	Variance of the estimated final residuals
<code>r2_qcorr</code>	R-square based on quadratic correlation between actual y and fitted y

Table 1: Bundle content as returned by `fsreg()`.

```
plot_rho_values(const bundle B)
```

Return type: void

#### 6.4 `plot_coefficient_paths()`

For plotting the development of the coefficients (coefficient paths), call the `plot_coefficient_paths()` function. It takes the resulting bundle returned by the `fsreg()` function. The argument is:

```
plot_coefficient_paths(const bundle B)
```

Return type: void



Parameter	Data type	Description	Default value
<code>eta</code>	scalar	Learning rate; $0 < \epsilon < 1$	0.05
<code>max_num_iterations</code>	int	Number of the maximum iterations	10000
<code>early_stopping_rounds</code>	int	Number of iterations of no improvement before stopping	50
<code>verbose</code>	bool	Print details or not: either 0 or 1	1 (True)

Table 2: Parameters which can be set through the optional bundle `opts`.